


Spark Streaming Data with Wings

Never stop the movement

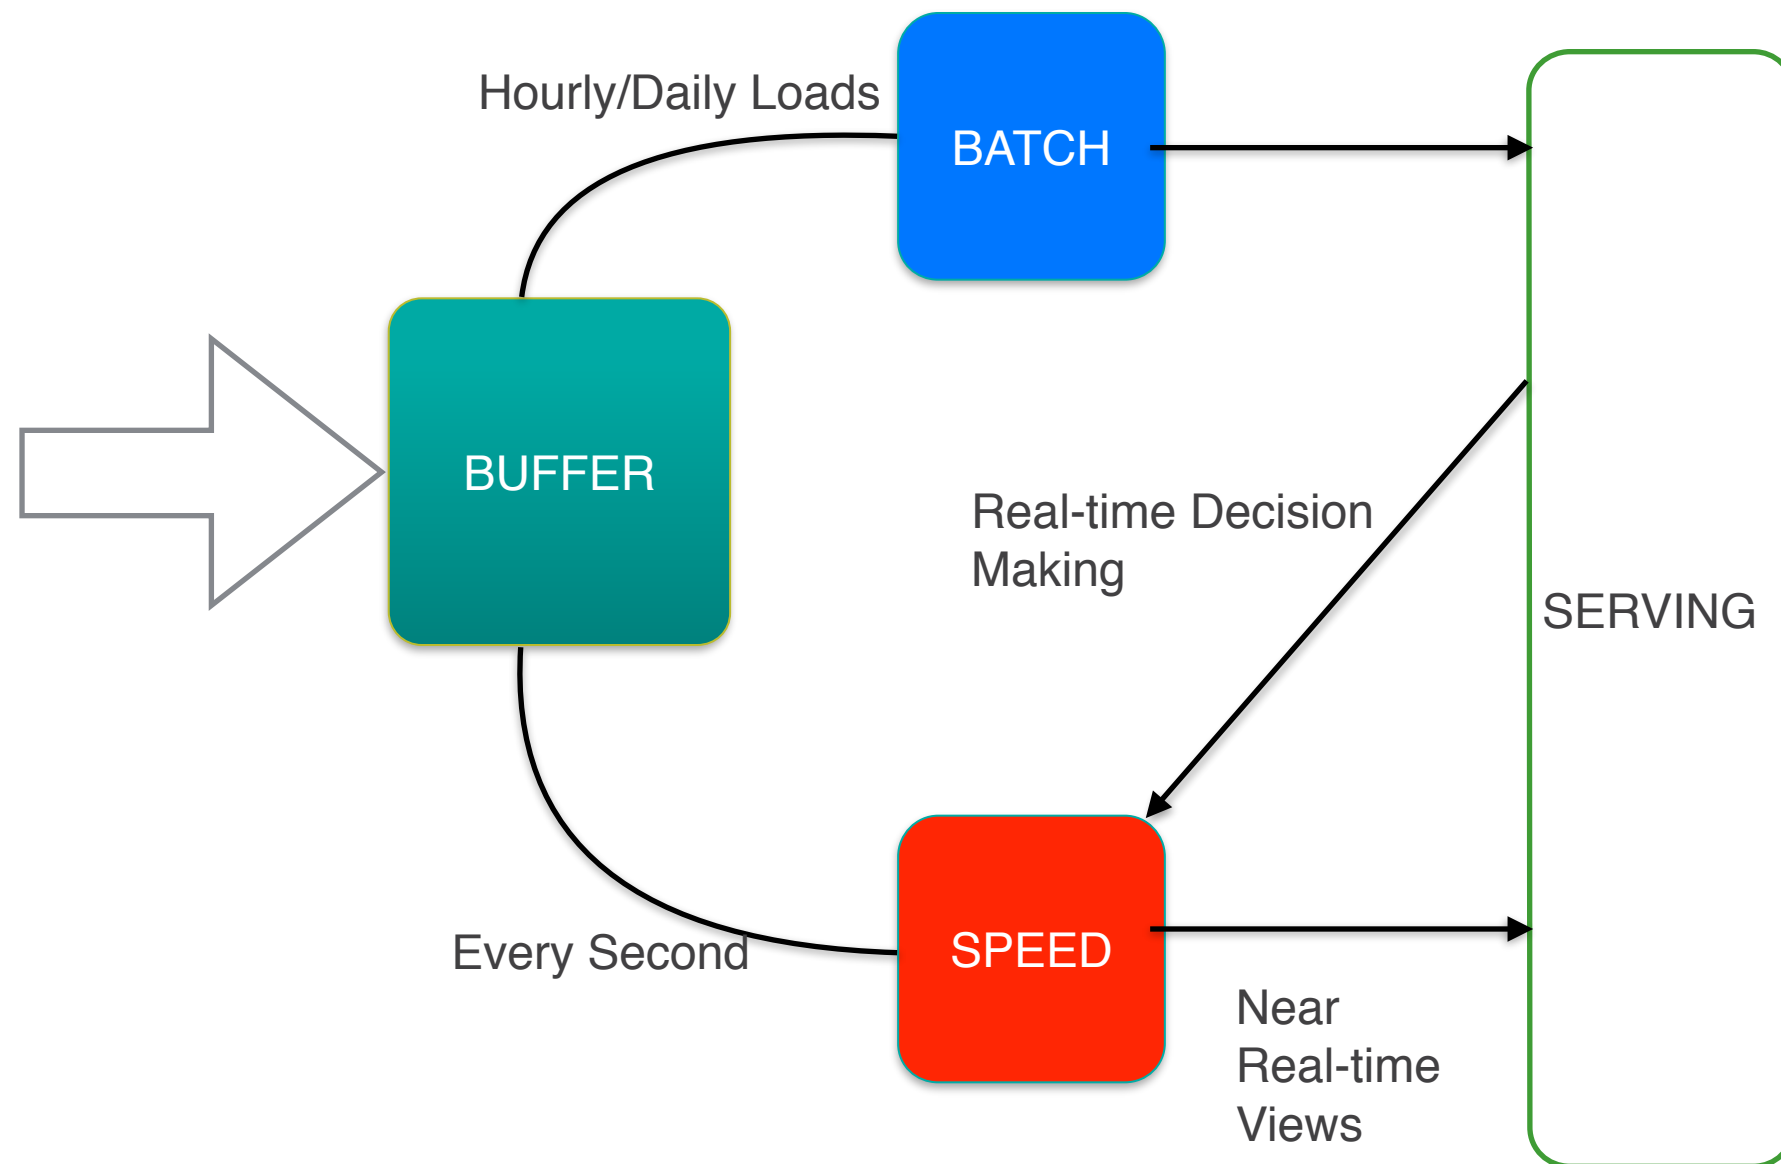
Agenda

- Spark overview
- Spark Streaming overview
- Processing streams from [meetup.com](https://www.meetup.com)
- Little intro to stateful stream processing
- All together

There will be no pictures of
cats!



Lambda Architecture



Do you see any problems
with Lambda Architecture?

Spark Overview

What if you could write programs operating with petabyte size **datasets** and large **streams** same way you operate Iterable collections?



Apache Spark

Spark
SQL

Spark
Streaming

MLlib
(machine
learning)

GraphX
(graph)

Apache Spark

Resilient Distributed Dataset

Resilient - resilience is addressed by tracking the log of operations performed on the dataset.

Because of the side effects are eliminated, every lost partition can be recalculated in case of a loss.

Distributed - the dataset is partitioned. We can specify partitioning scheme for every operation.


Dataset - can be built from regular files, HDFS large files, Cassandra table, HBase table, etc.

Obligatory word count

```
lines = spark.textFile("hdfs://...")
```

```
lines.flatMap{line: line.split(" ")}  
  
    .map({word: (word, 1)})  
  
    .reduceByKey(_+_)
```

Spark Runtime




Spark is like Yoda and Hulk
combined



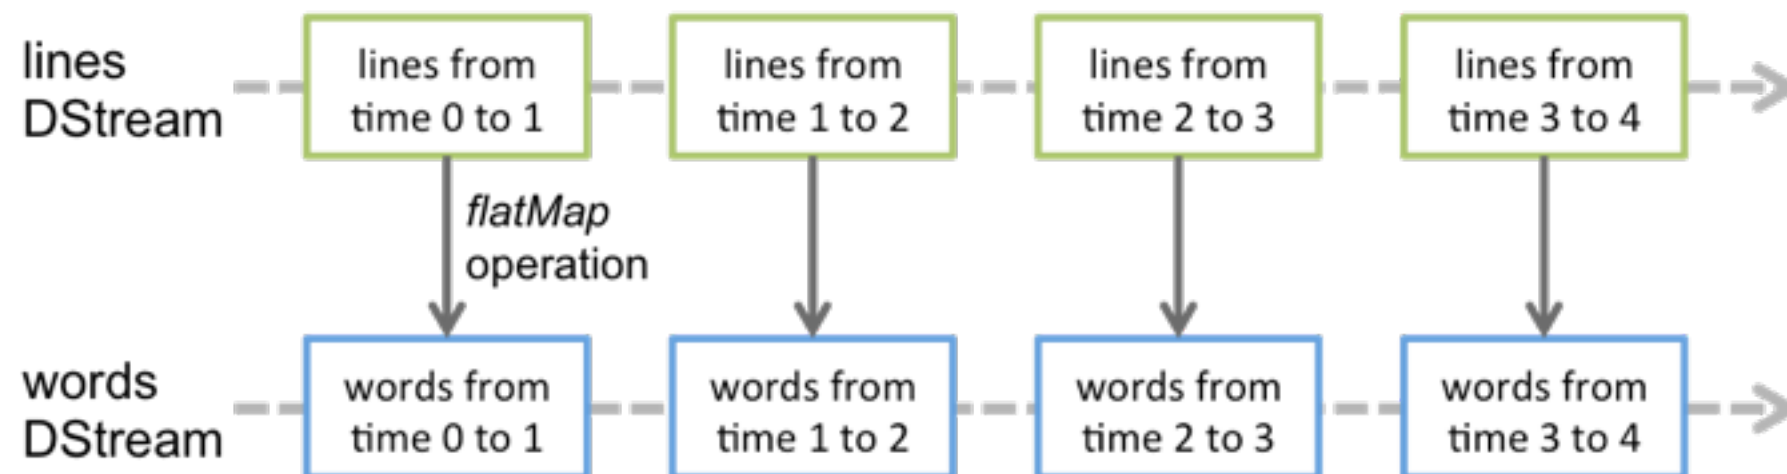
Discretized Stream

New RDD every second



DStream - still pretty

You operate on stream just like on collections



<https://spark.apache.org/docs/1.2.0/streaming-programming-guide.html>

meetup.com Streams

Let's play with meetup.com streams:

- Events
- RSVPs

RSVP Schema

```
{ "event" : { "event_id" : "220993343",  
  "event_name" : "Paper Collage",  
  "event_url" : "http://www.meetup.com/ELEOS-Art-School-Studio/events/220993343/",  
  "time" : 1427551200000  
},  
  "group" : {...},  
  "guests" : 0,  
  "member" : { "member_id" : 120762942,  
    "member_name" : "Esther",  
    "photo" : "http://photos1.meetupstatic.com/photos/member/b/b/0/thumb_159962992.jpeg"  
  },  
  ...  
}
```

Event Schema

```
{ "description" : "<p>90 Minute walking tour with your dog! Please arrive early.</p>\n<p>Tour of Balboa Parks spookiest locations on ...>",  
  "duration" : 5400000,  
  "event_url" : "http://www.meetup.com/SanDiegoDogWalkers/events/220302036/",  
  "group" : {...},  
  "id" : "220302036",  
  "maybe_rsvp_count" : 0,  
  "mtime" : 1425785739616,  
  "name" : "After Dark Ghost Walking Tour in Balboa Park with your dog!",  
  "payment_required" : "0",  
  "rsvp_limit" : 20,  
  "status" : "upcoming",  
  "time" : 1426993200000,  
  "utc_offset" : -25200000,  
  "venue" : {...},
```


Meetup Receiver

Receiver is the way to pump data to spark streaming.

In our case, we connect to meetup streaming api and send json to Spark.

Code is a bit long, but you can explore it:

<https://github.com/actions/meetup-stream/blob/master/src/main/scala/receiver/MeetupReceiver.scala>

Intro to Stateful Stream Processing

```
def locaitonCounts(eventsStream: DStream[(Venue, Event)]) =  
  liveEvents  
    .filter{case(venue, event)=>venue.country=="usa"}  
    .map{case(venue, event)=>(venue.toCityState, 1)}  
    .updateStateByKey(countForLocaiton)  
    .print
```

```
def countForLocaiton(counts: Seq[Int], initCount: Option[Int])  
=Some(initCount.getOrElse(0) + counts.sum)
```

How it works again...?

Incoming Stream

CityState	Count
San Francisco,CA	1
Miami, FL	1

CityState	Count
San Francisco,CA	1
Miami, FL	1

10 sec → 11 sec

CityState	Count
San Francisco, CA	5
LA, CA	3
Portland, OR	2

CityState	Count
San Francisco, CA	6
LA, CA	3
Portland, OR	2
Miami, FL	1

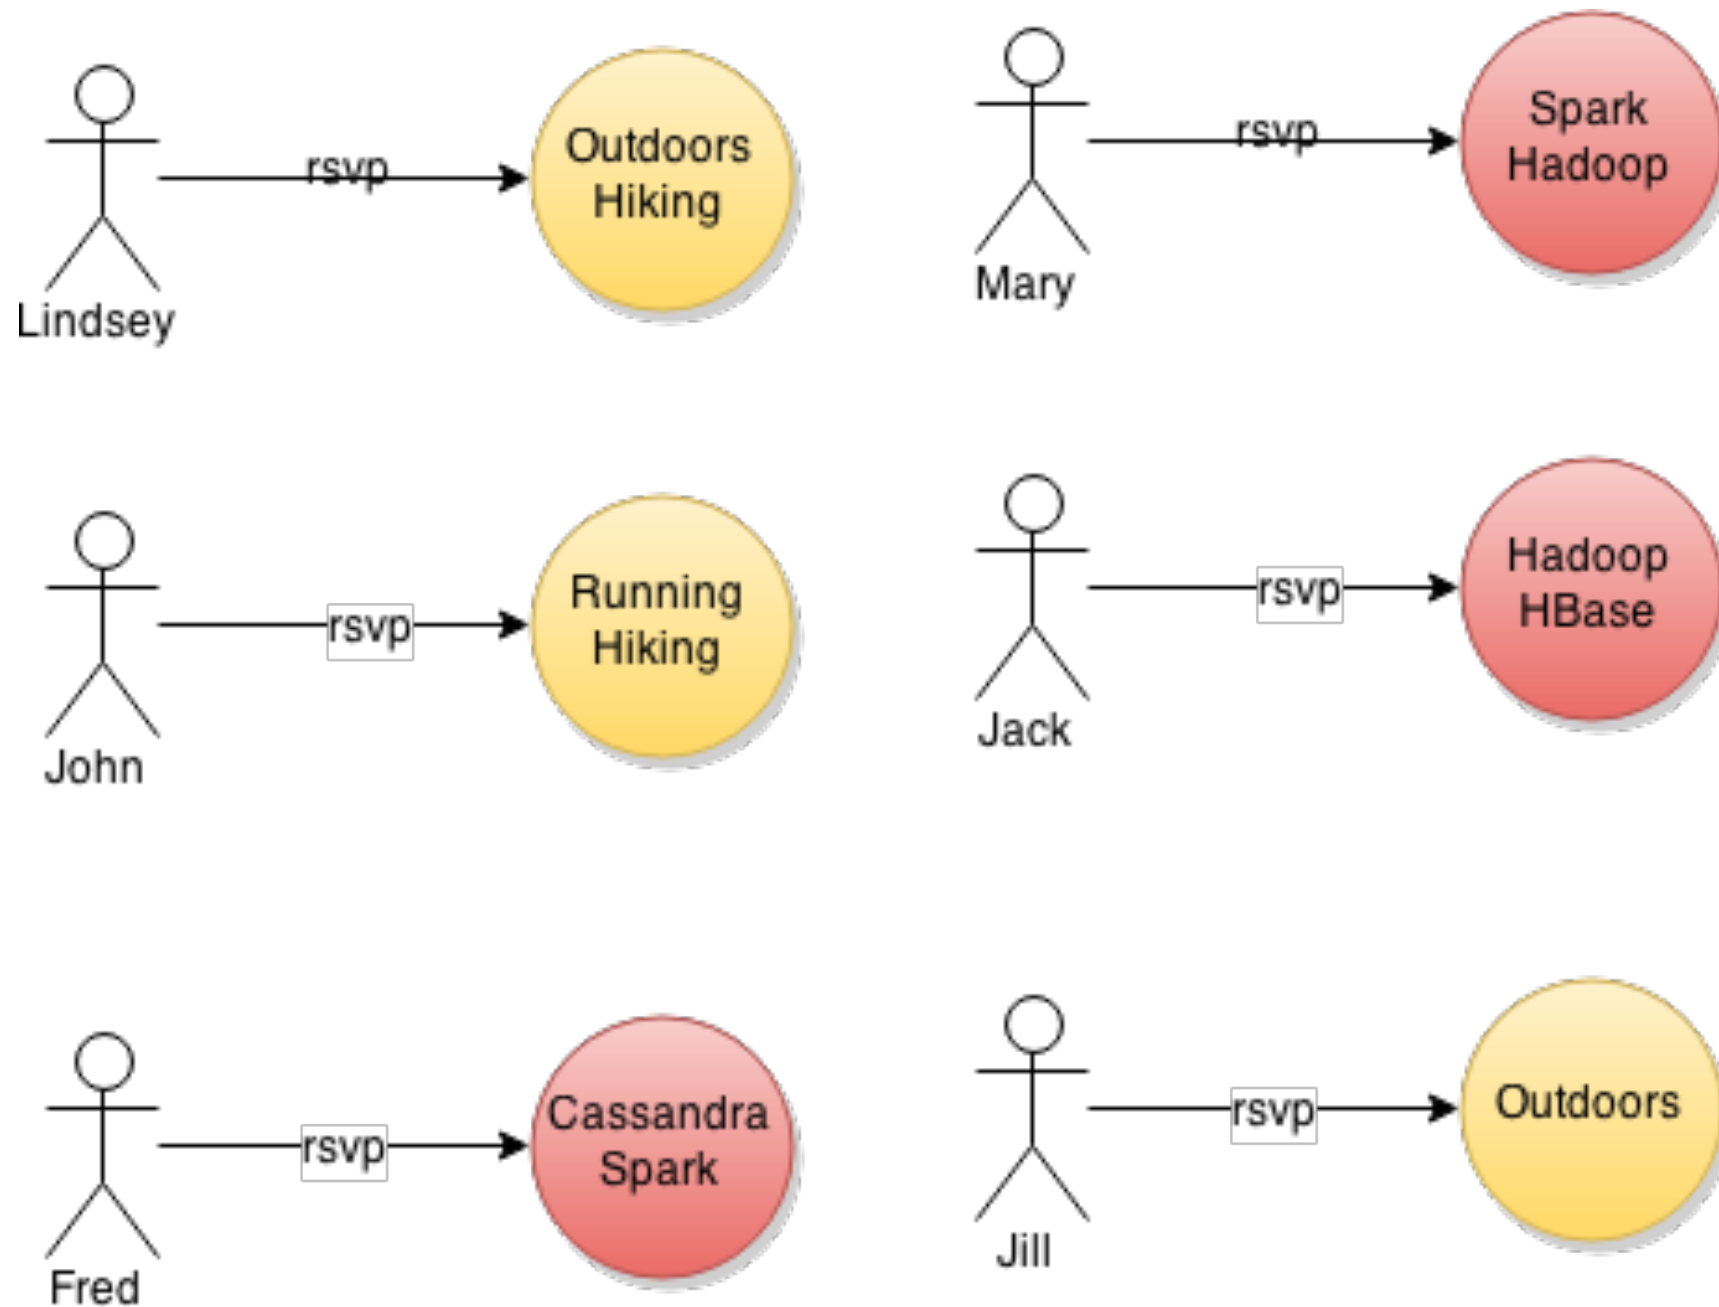
State Stream

Aerial refueling

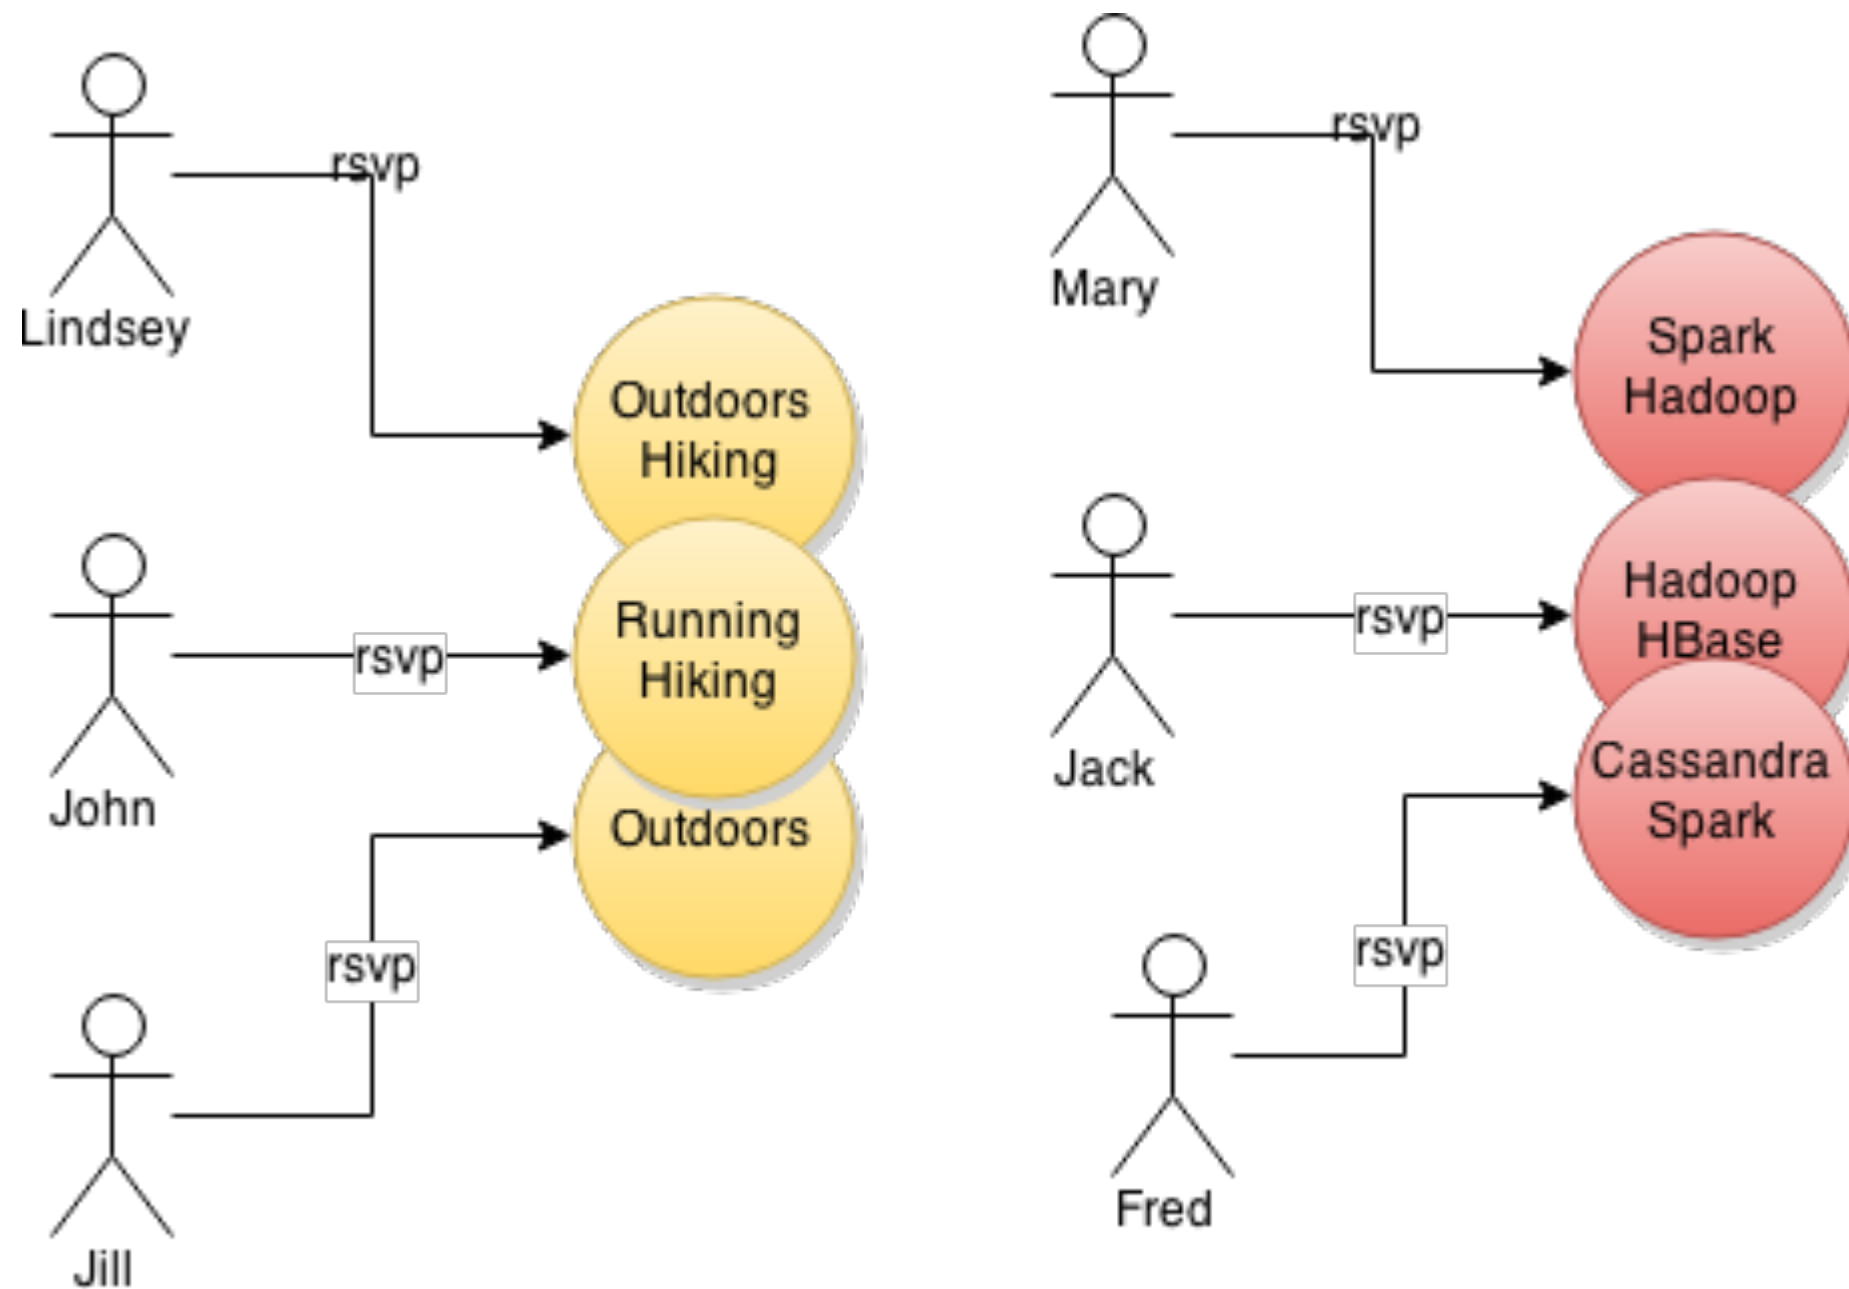


http://commons.wikimedia.org/wiki/File:Aerial_refueling_CH-53_DF-SD-06-02984.JPG



meetup.com connection recommendation app



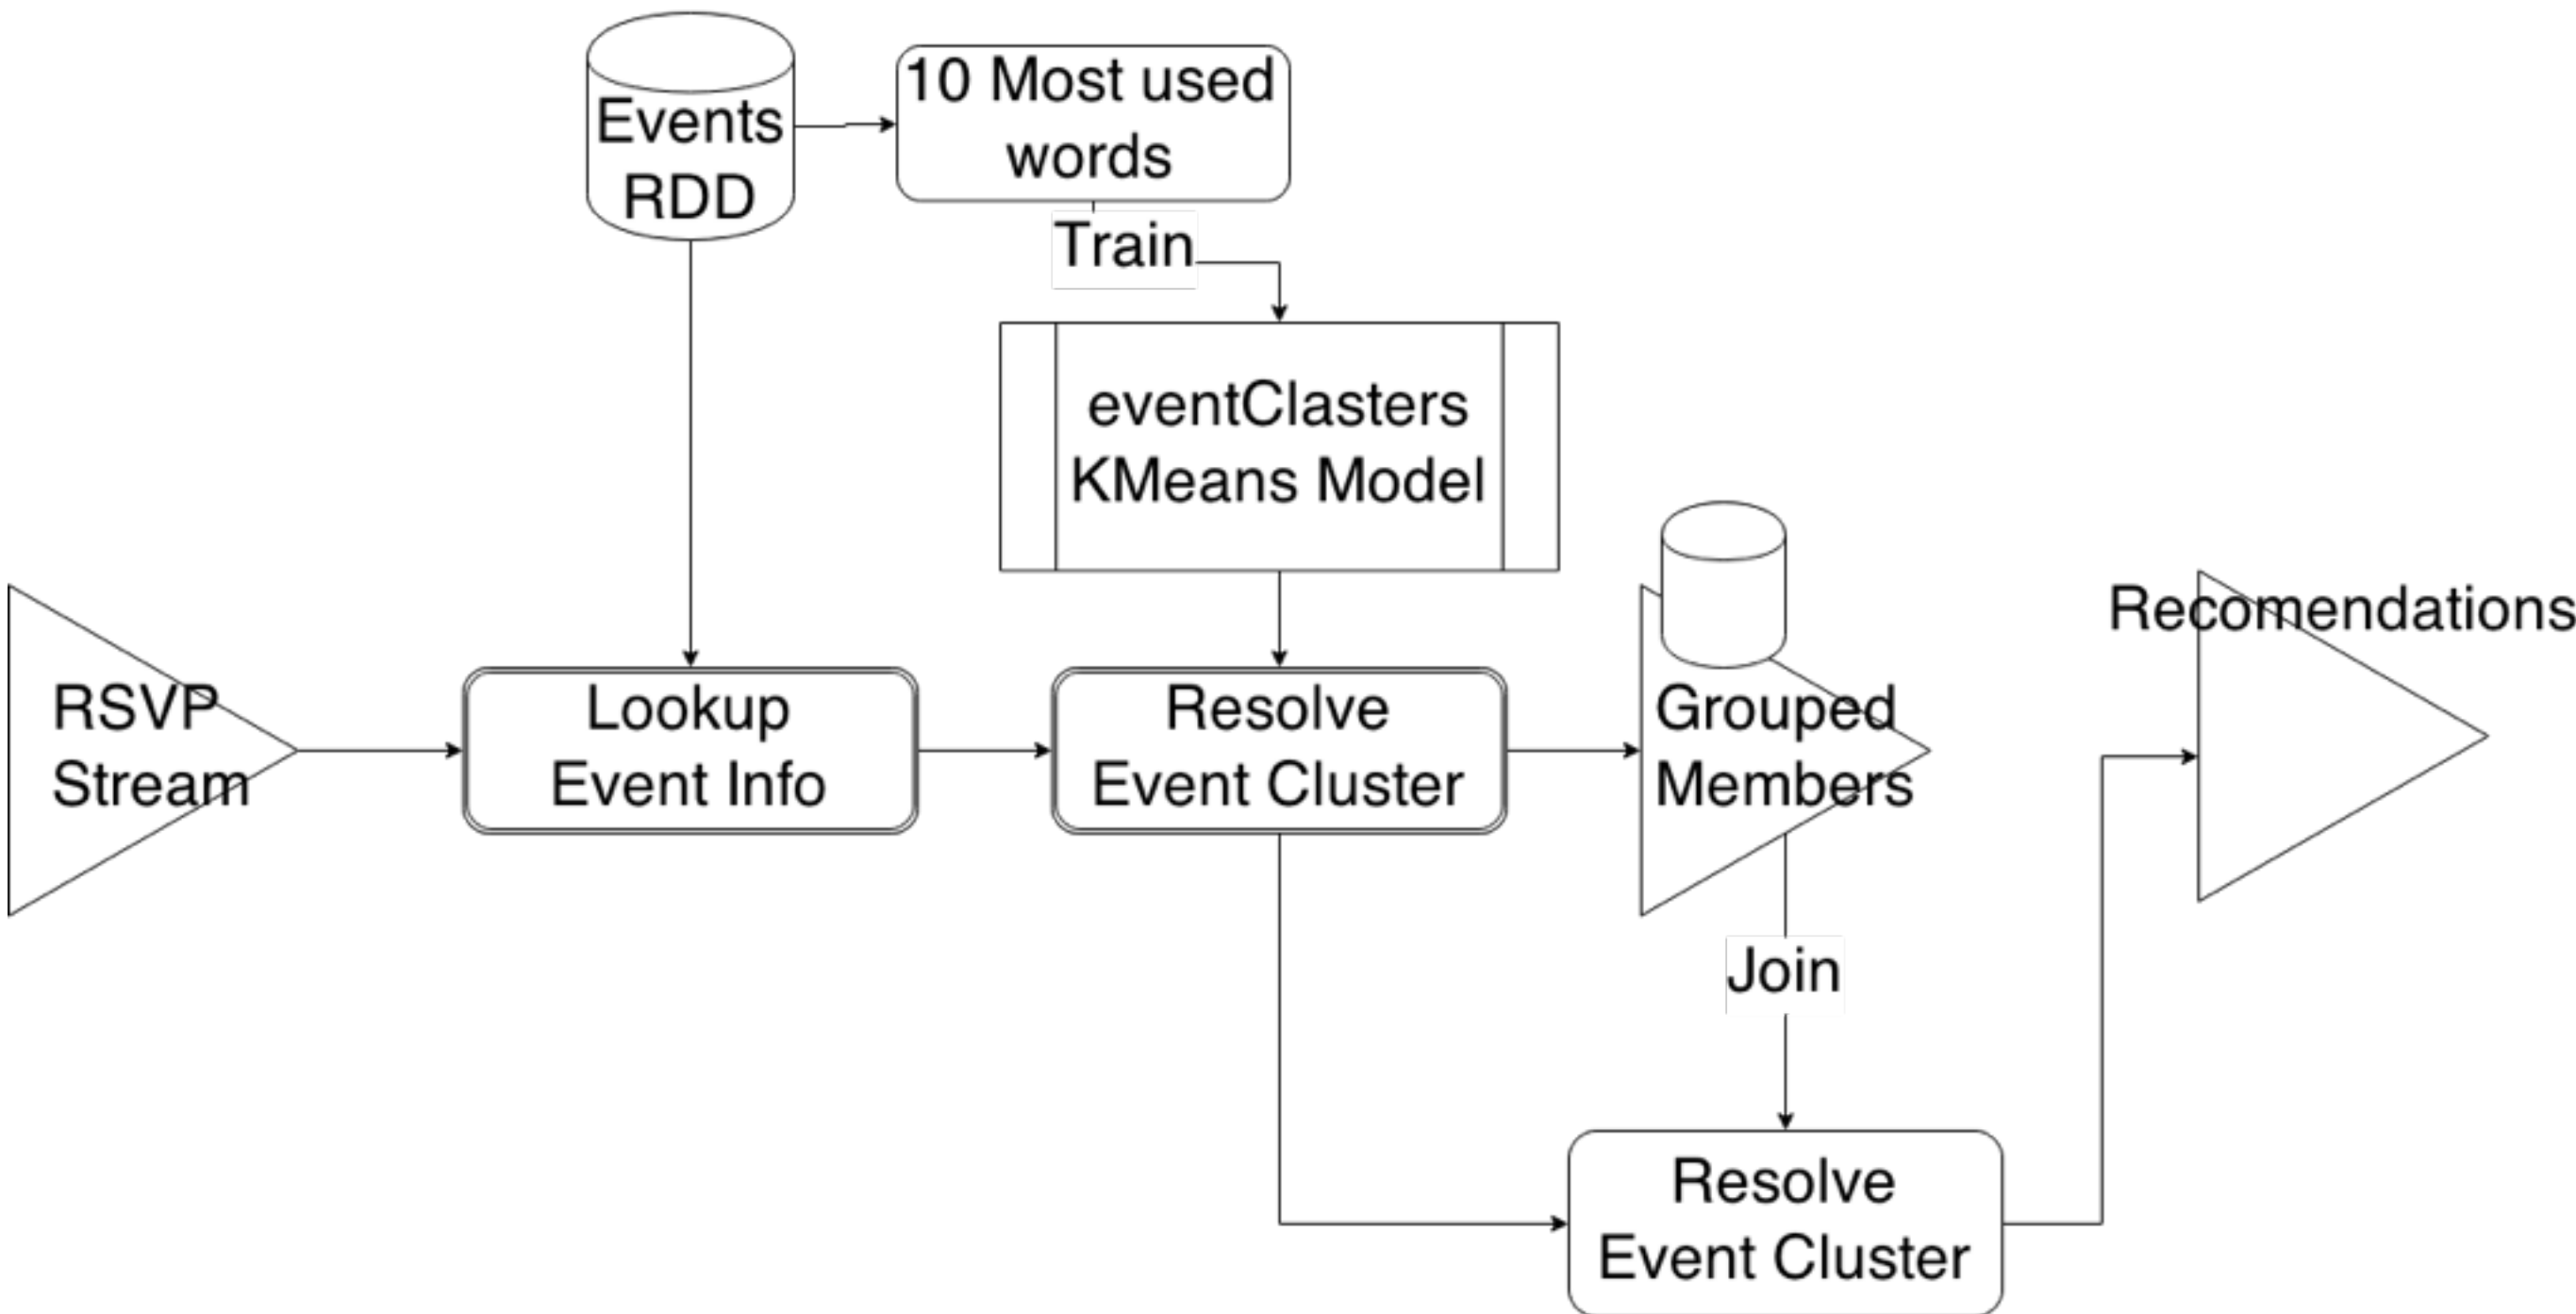
Meetup Event Clustering



Meetup Professional Connection Recommendation



Meetup Recommendations Pipeline



Initializing RSVP Stream and the Event Dataset

```
val conf = new SparkConf()
    .setMaster("local[4]")
    .setAppName("MeetupExperiments")
    .set("spark.executor.memory", "1g")
    .set("spark.driver.memory", "1g")
```

```
val ssc=new StreamingContext(conf, Seconds(1))
```

```
val rsvpStream = ssc.receiverStream(
    new MeetupReceiver("http://stream.meetup.com/2/rsvps")).flatMap(parseRsvp)
```

```
val eventsHistory = ssc.sparkContext.textFile("data/events/
events.json", 1).flatMap(parseEvent)
```

Broadcasting Dictionary

```
val localDictionary=Source
    .fromURL(getClass.getResource("/wordsEn.txt"))
    .getLines
    .zipWithIndex
    .toMap

val dictionary=ssc.sparkContext
    .broadcast(localDictionary)
```


Feature Extraction

10 most popular words in the description.

```
def eventToVector(event: Event): Option[Vector]={  
    val wordsIterator =  
event.description.map(breakToWords).getOrElse(Iterator())  
    val topWords=popularWords(wordsIterator)  
    if (topWords.size==10)  
Some(Vectors.sparse(dictionary.value.size,topWords)) else None  
}  
  
val eventVectors=eventsHistory.flatMap{event=>eventToVector(event)}
```

Training based on existing dataset

K-means clustering on the digits dataset (PCA-reduced data)
Centroids are marked with white cross



```
val eventClusters = KMeans.train(eventVectors, 10, 2)
```

http://scikit-learn.org/0.11/auto_examples/cluster/plot_kmeans_digits.html

Event History By ID

```
val eventHistoryById=eventsHistory  
    .map{event=>(event.id, event.description.getOrElse(""))}  
    .reduceByKey{(first: String, second: String)=>first}
```

```
(220302036, "...description1 ...")  
(220302037, "...description2 ...")  
(220302038, "...description3 ...")
```

Streaming lookups

Looking up the event description by eventId from rsvp.

```
val rsvpEventInfo = membersByEventId.transform(  
  rdd=>rdd.join(eventHistoryById)  
)
```

(eventId, (member, response), description)

(220819928, ((Member(Some(cecilia rogers), Some(162556712)), yes), "...")

(221153676, ((Member(Some(Carol), Some(183499291)), no), "...")

...

Streaming Clustering

```
val memberEventInfo = rsvpEventInfo
  .flatMap{
    case(eventId, ((member, response), event)) => {
      eventToVector(event).map{ eventVector=>
        val eventCluster=eventClusters.predict(eventVector)
        (eventCluster,(member, response))
      }
    }
  }
```

Clustering members

```
def groupMembers(memberResponses: Seq[(Member, String)], initList: Option[Set[Member]]) =  
{  
    val initialMemberList=initList.getOrElse(Set())  
    val newMemberList=(memberResponses :\ initialMemberList) {  
        case((member, response), memberList) =>  
            if (response == "yes") memberList + member else memberList - member  
    }  
    if (newMemberList.size>0) Some(newMemberList) else None  
}  
  
val memberGroups =  
memberEventInfo.updateStateByKey(groupMembers)
```


Recommendations

```
val recommendations=memberEventInfo
    .join(memberGroups)
    .map{
        case(cluster, ((member, memberResponse), members)) =>
            (member.memberName, members-member)
    }
```

```
(Some(Mike D) -> Set(Member(Some(Sioux),Some(85761302)),
Member(Some(Aileen),Some(12579632)),
Member(Some(Teri),Some(148306762))))
```

Questions

